

Linker Relocation

Peter Smith

Introduction and purpose

- Why do object files have relocations?
- How is a relocation represented in an Object file?
- How does a linker resolve a relocation?
- How to read the ABI documentation?
- GOT and PLT generating relocations and dynamic relocations?
- Relative relocations

Why do we have relocations?

- Linking is the process of turning the abstract into the concrete
- Compilers and assemblers use symbolic references
- Assembler can resolve symbolic “fixups” within the same section
- Static linker can resolve offsets between sections and knows absolute addresses
- Dynamic linker can resolve references between shared libraries and relative displacements
- Relocations express the calculation that the linker or loader must perform

How is a relocation represented in an object file

```
int val;
```

```
int func(void) {  
    return val;  
}
```

```
clang --target=aarch64-linux-gnu -O2
```

```
.global val
```

```
0000000000000000 adrp x8, val  
// 0000000000000000: R_AARCH64_ADR_PREL_PG_HI21      val  
0000000000000004 ldr w0, [x8, :lo12:val]  
// 0000000000000004: R_AARCH64_LDST32_ABS_LO12_NC      val  
0000000000000008 ret
```

Differences when compiling with PIC

```
int val; // Preemptible
```

```
int func(void) {  
    return val;  
}
```

```
clang --target=aarch64-linux-gnu -O2 -fpic
```

```
.global val
```

```
0000000000000000 adrp x8, :got: val  
// 0000000000000000: R_AARCH64_ADR_GOT_PAGE val  
0000000000000004 ldr x8, [x8, :got_lo12:val]  
// 0000000000000004: R_AARCH64_LD64_GOT_LO12_NC val  
0000000000000008 ldr w0, [x8]  
000000000000000c ret
```

How does a linker resolve a relocation

- Relocation resolution is the last action before writing the file
 - Image layout must be complete
- Relocation code identifies all the actions a linker needs to take
 - Smart format, dumb linker
- Simplified steps, with S as VA of symbol, P as VA of place of relocation
 - Extract the addend A from the instruction (REL only)
 - Perform the calculation, most common ones are $S + A$ (absolute) or $S + A - P$ (relative)
 - Perform any alignment or overflow checks
 - Encode the result of the calculation in the instruction

How to read the ABI documentation

- Will also need the Architecture Reference manual for the instruction encodings
- Descriptions of relocation operators used in the tables
 - “**G DAT (S+A)** represents a pointer-sized entry in the **GOT** for address **S+A**. The entry will be relocated at run time with relocation **R_<CLS>_GLOB_DAT (S+A)**.”
- Tables of relocation codes, sometimes partitioned into related relocations
 - Almost always partitioned into static and dynamic
 - “Table 36 Table 4-14, GOT-relative instruction relocations”
- At a minimum each relocation has code, calculation, remarks
 - Fill in the gaps with the Architecture reference manual and some imagination
- Relocations with special behavior such as PLT, GOT generation

Example R_AARCH64_ADR_PREL21

ELF for the 64-bit Arm Architecture table entry

275	11	R <CLS> ADR_PREL_PG_HI21	Page (S+A) - Page (P)	Set an ADRP immediate value to bits [32:12] of the X; check that $-2^{32} \leq x < 2^{32}$
-----	----	-----------------------------	-----------------------	--

ADRP instruction Arm ARM entry


1	Immlo (2)	1 0 0 0 0	Immhi (18)	Rd (5)
---	-----------	-----------	------------	--------

`imm = SignExtend(immhi:immlo:Zeros(12), 64);`

Linker implementation

- `imm = checkIntOverflow(((S + A) & ~0xfff) - (P & ~0xfff)), 33) >> 12`
- `immLo = (imm & 0x3) << 29 ; immHi = (imm & 0x1fffffc << 3);`
- `immMask = (0x3 << 29) | (0x1fffffc << 3)`
- `write32le (buf, (read32le(buf) & ~immMask) | immLo | immHi);`

Visualisation



R_AARCH64_ADR_PREL21 :

```
Page (S + A) - Page (P) =  
0x100000 - 0x8000 = 0xf8000  
adrp x8, 0xf800  
x8 = 0x8000 + f8000 = 0x10000
```

R_AARCH64_LDST32_ABS_LO12_NC

```
(S + A) & 0xffc  
0x100040 & 0xffc = 0x40  
ldr w0, [x8, #40]
```

GOT and PLT generating relocations

.text
.plt
.rodata

Read-only
no relocations
shareable

.data.rel.ro
.got

Read-only after
relocation
not-shareable

.got.plt
.data

Read-write
not-shareable

Procedure Linkage Table **PLT**

- Linker generated stubs to call functions in shared libraries
- Ifuncs
- Load addresses from the .got.plt
- Lazy binding makes .got.plt RW
- Entry generated by call or jump relocation to function that could be defined externally
- .rela.plt

Global Offset Table **GOT**

- Pointers to data
- Resolved at load time
- Can be made read-only after relocation
- Relocations can be simplified to `R_<arch>_RELATIVE`
- .rela.dyn


GOT generating relocation example

ELF for the 64-bit Arm Architecture entry

311	16	R <CLS> ADR_PREL_PG_HI21	Page (G (GDAT (S+A))) - Page (P)	Set an ADRP immediate value to bits [32:12] of the X; check that $-2^{32} \leq X < 2^{32}$
-----	----	-----------------------------	---	--

- **GDAT (S+A)** represents a pointer-sized entry in the GOT for address **S+A**. The entry will be relocated at run time with relocation **R_<CLS>_GLOB_DAT (S+A)**.
- **G (expr)** is the address of the GOT entry for the expression **expr**.
- Linker creates GOT slot for **val**, if it doesn't already exist
- Linker creates **R_AARCH64_GLOB_DAT** relocation for got slot with symbol **val**
- Linker evaluates **Page(Address of got slot) - Page (position)**
- Linker writes back the offset to the page that the GOT slot is in.
- The **R_AARCH64_LD64_GOT_LO12_NC** loads the low 12 bits as in previous example

Visualisation



R_AARCH64_ADR_GOT_PAGE :
Page G(GDAT(S + A)) - Page (P)
=
0x9000 - 0x8000 = 0x1000
adrp x8, 0x1000
x8 = 0x8000 + 0x1000 = 0x9000

R_AARCH64_LD64_GOT_LO12_NC
G(GDAT(S + A)) & 0xffc
0x9040 & 0xffc = 0x40
ldr x8, [x8, #40]

Relative relocations

- The R_<Arch>_GLOB_DAT dynamic relocation is expensive to resolve
 - Dynamic loader has to lookup symbol by name
- When the definition of a symbol is DSO Local then the R_<Arch>_RELATIVE dynamic relocation can be used
 - Linker writes the value of the address at static link time into the GOT slot
 - Relative relocation is the displacement from the static link address to the runtime address
 - A simple addition, no symbol lookups, all relative relocations get the same value
 - Sorted before symbol lookups for ease of processing by the dynamic loader
- PIE binaries often have large numbers of relative relocations
- Android compresses relative relocations

Further topics not covered today

- Lazy binding via PLT sequences
- Thread local storage models and relaxation
- Linker generated thunks
- Garbage collection and identical code folding
- String and constant merging
- Comdat groups
- Linker Scripts
- Symbol ordering files